

# **DEEP LEARNING**

Lecture

Introduction

*Instructor: Zafar Iqbal*

# Agenda

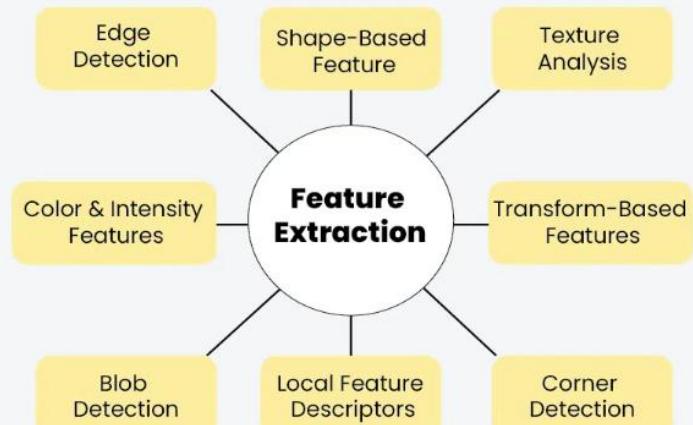
- What is Deep Learning?
- Difference between AI, ML and DL
- History, Key components of DL
- Structure, Applications, Challenges
- Frameworks, Mathematics



# What is Deep Learning?

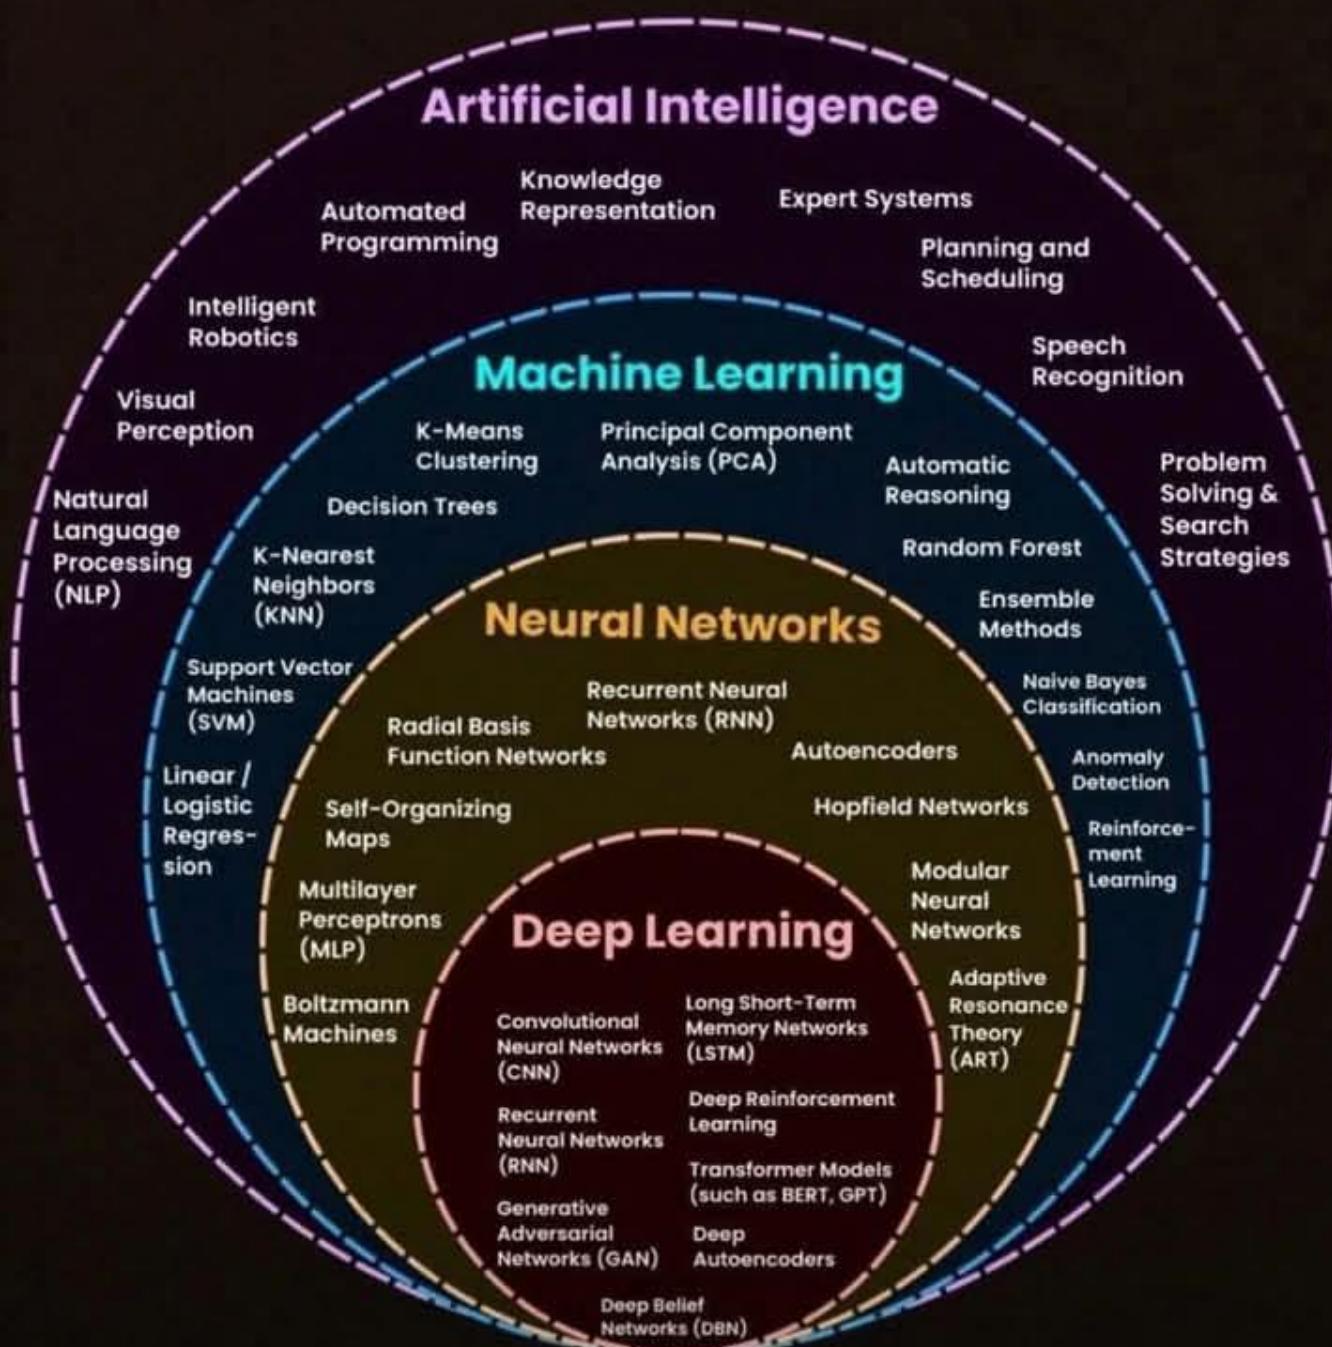
- A subset of machine learning based on artificial neural networks
- Inspired by the structure and function of the human brain
- Learns hierarchical data representations
- Automatically extracts features from raw data
- Enables machines to perform tasks like humans (vision, language, speech)

## Feature Extraction Techniques in Image Processing



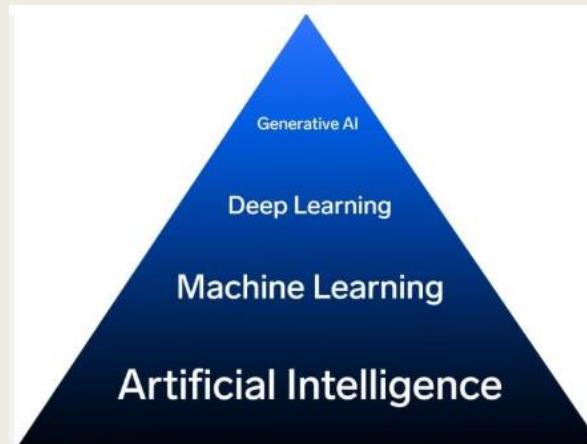
# What is Deep Learning

- Neural networks have become the defining model of deep learning due to their power and scalability.
- They are composed of neurons, each of which performs a simple computation.
- The true power of a neural network comes from the complexity of the connections between these neurons.
- These complex connections enable the network to model complicated patterns and relationships in data.



# AI vs ML vs DL

Level	Description	Example
AI	Broad goal: make machines intelligent	Chatbots, self-driving cars
ML	Algorithms that learn from data	Decision Trees, SVMs
DL	Neural networks with multiple layers	CNNs, RNNs, Transformers



# Why Deep Learning?

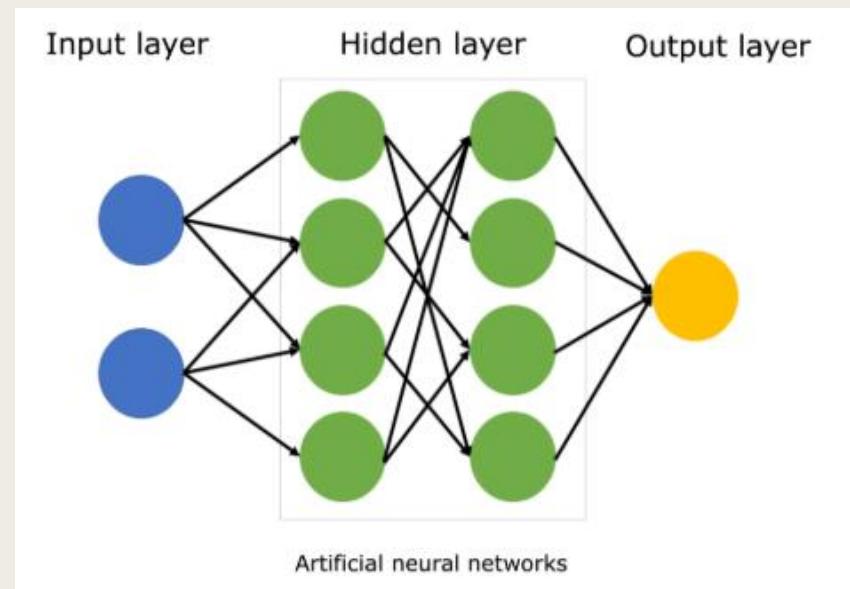
- Handles complex, high-dimensional data
- Requires less manual feature engineering
- Scales well with big data and GPUs
- Achieves state-of-the-art performance in multiple domains

# History of deep learning

- 1943: McCulloch-Pitts neuron model
- 1986: Backpropagation algorithm introduced
- 2006: Deep Belief Networks by Hinton et al.
- 2012: AlexNet wins ImageNet, sparking deep learning boom
- 2020s: Transformers and large-scale models dominate

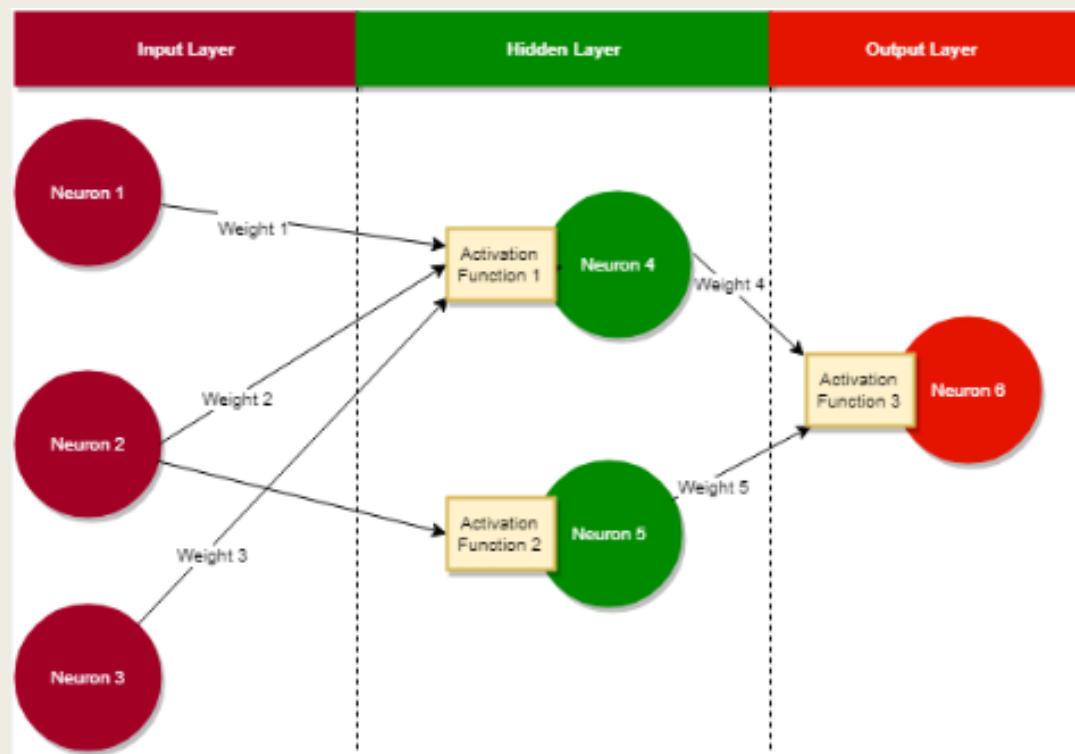
# Key Components of Deep Learning

- **Neural Networks:** Core computation structure
- **Data:** Large datasets for learning patterns
- **Computation Power:** GPUs/TPUs for acceleration
- **Algorithms:** Optimization (SGD, Adam), activation functions, loss functions



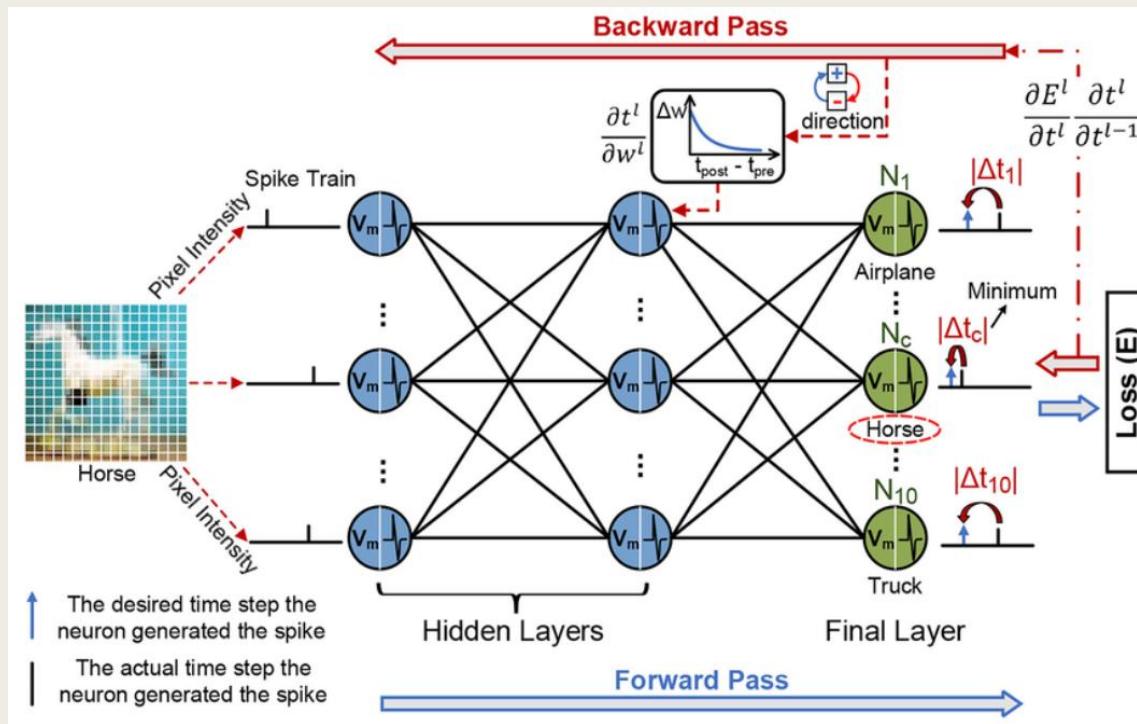
# Structure of a Neural Network

- Input layer
- Hidden layers
- Output layer
- Weights and biases
- Activation functions



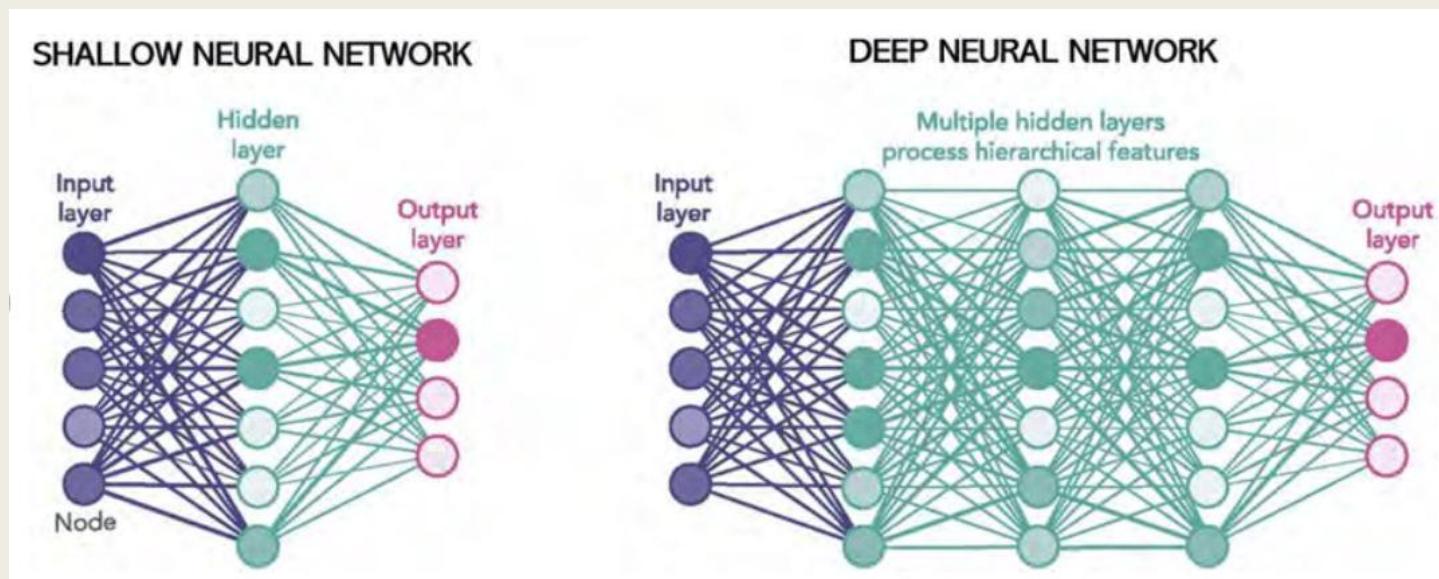
# How Neural Networks Learn

- Forward pass: computes predictions
- Loss function: measures prediction error
- Backpropagation: updates weights via gradients
- Optimization loop continues until convergence



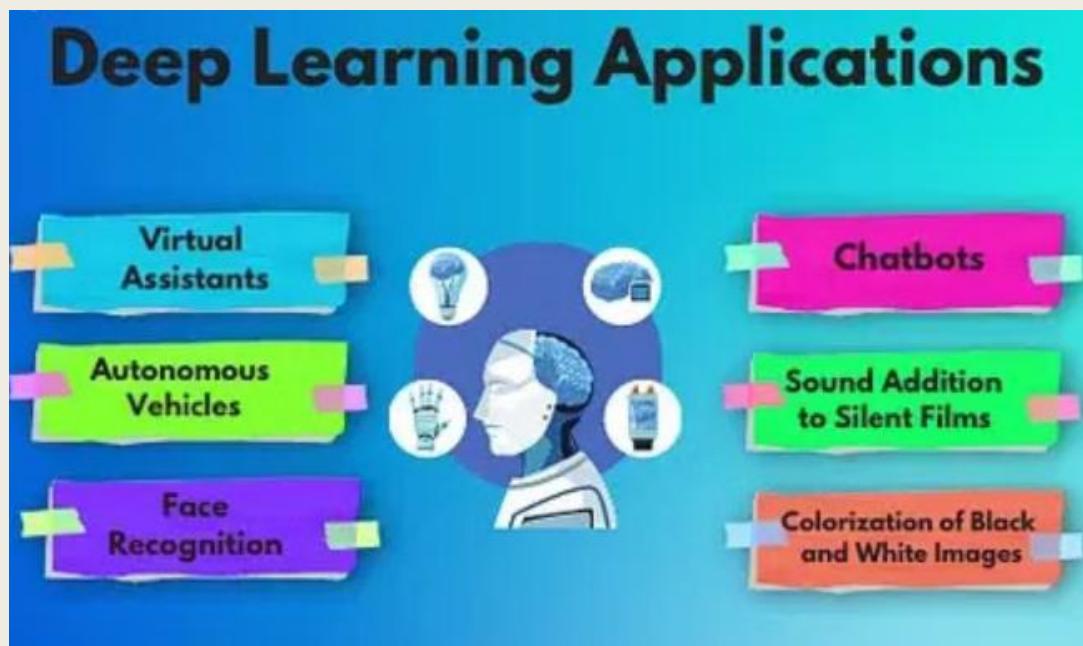
# Deep vs Shallow Networks

Feature	Shallow	Deep
Layers	1–2 hidden layers	Many layers
Features	Hand-crafted	Learned automatically
Data Need	Small	Large
Accuracy	Limited	High (if data-rich)



# Common Applications

- Computer Vision: Object detection, facial recognition
- Natural Language Processing: Chatbots, translation
- Speech Recognition: Voice assistants
- Healthcare: Disease diagnosis, image analysis
- Autonomous Systems: Self-driving cars, robotics



# Real-World Examples

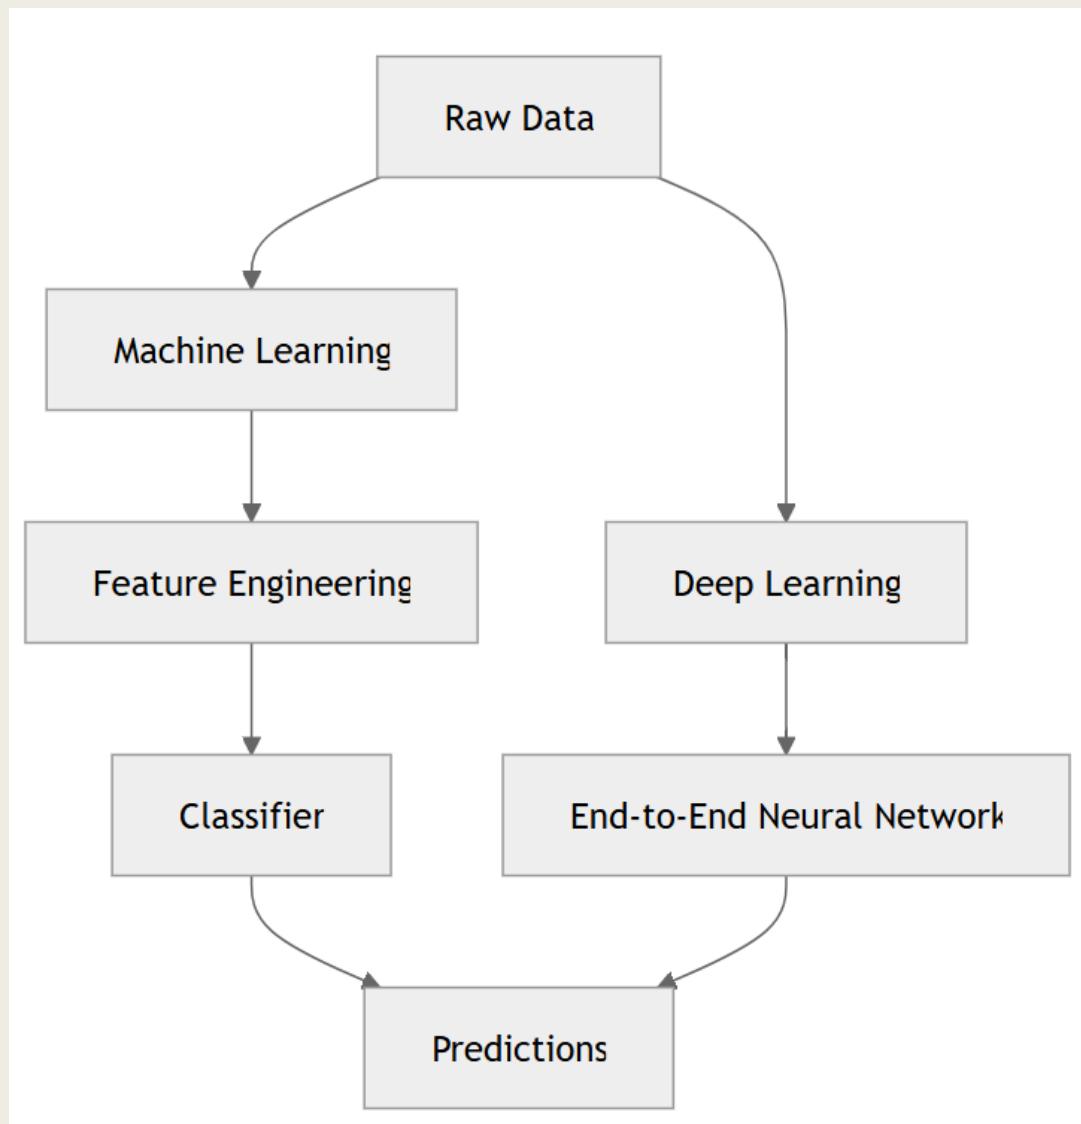
- Google Translate (Sequence-to-sequence models)
- ChatGPT (Transformer models)
- Tesla Autopilot (CNNs for vision)
- DeepMind AlphaGo (Reinforcement learning)



# Challenges in Deep Learning

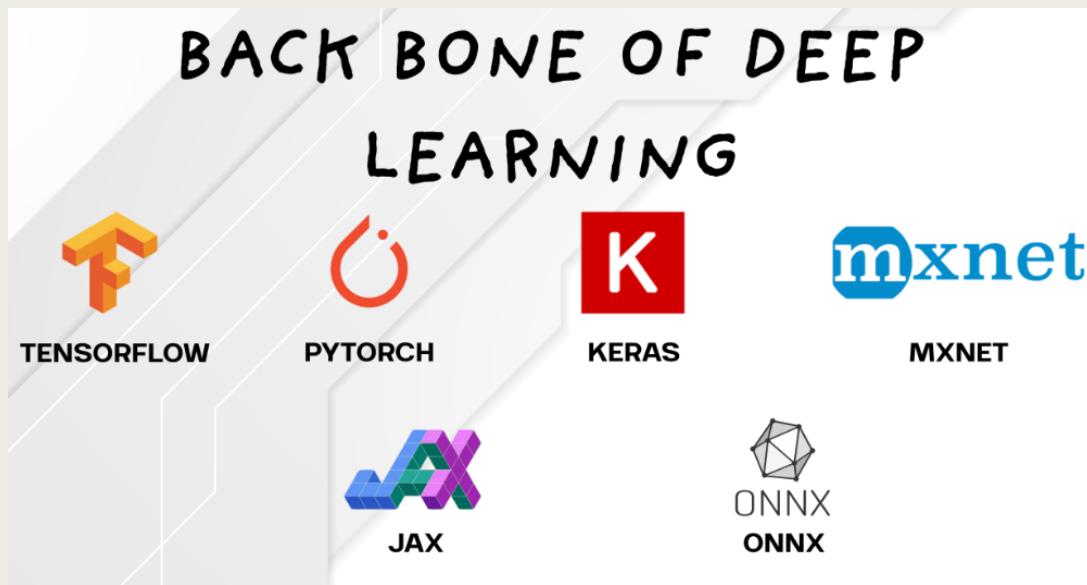
- Data requirements (large labeled datasets)
- Computational cost (training time & GPU needs)
- Interpretability (black-box nature)
- Bias and fairness issues
- Overfitting

# Machine Learning vs Deep Learning



# Popular Deep Learning Frameworks

- TensorFlow (Google)
- PyTorch (Meta)
- Keras (High-level API)
- JAX, MXNet, ONNX  
Open-source, GPU-accelerated, and widely supported



# Typical Deep Learning Workflow

- Data collection and preprocessing
- Model design (architecture)
- Training (optimize weights)
- Evaluation (metrics, validation)
- Deployment and monitoring

# Deep Learning Ecosystem

- **Hardware:** GPUs, TPUs
- **Software:** Libraries and frameworks
- **Data:** Public datasets (ImageNet, COCO, etc.)
- **Community:** Research papers, GitHub repos

# Why Use GPUs?

- Parallel processing for matrix operations
- Speeds up training by 10–100x compared to CPUs
- Cloud options: AWS, Google Cloud, Azure [Image: GPU vs. CPU performance comparison]



# Linear Algebra Basics

- Vectors: Represent data points or weights
- Matrices: Represent transformations or connections
- Tensors: Generalize vectors/matrices to higher dimensions
- Operations: Dot product, matrix multiplication Equation:  
 $y = Wx + b$

# Calculus for Optimization

- Gradients: Measure rate of change of loss function
- Partial Derivatives: Optimize weights in neural networks
- Gradient Descent: Minimize loss iteratively Equation:
- $w \leftarrow w - \eta * \partial L / \partial w$

# Probability and Statistics

- Probability: Model uncertainty (e.g., softmax outputs)
- Expectations: Used in loss functions
- Distributions: Gaussian, Bernoulli for data modeling  
Equation:  
$$P(y|x) = \exp(w_y * x) / \sum(\exp(w_i * x))$$
 (Softmax)