

# **DEEP LEARNING**

## Lecture 2

Simple Neural Network

*Instructor: Zafar Iqbal*

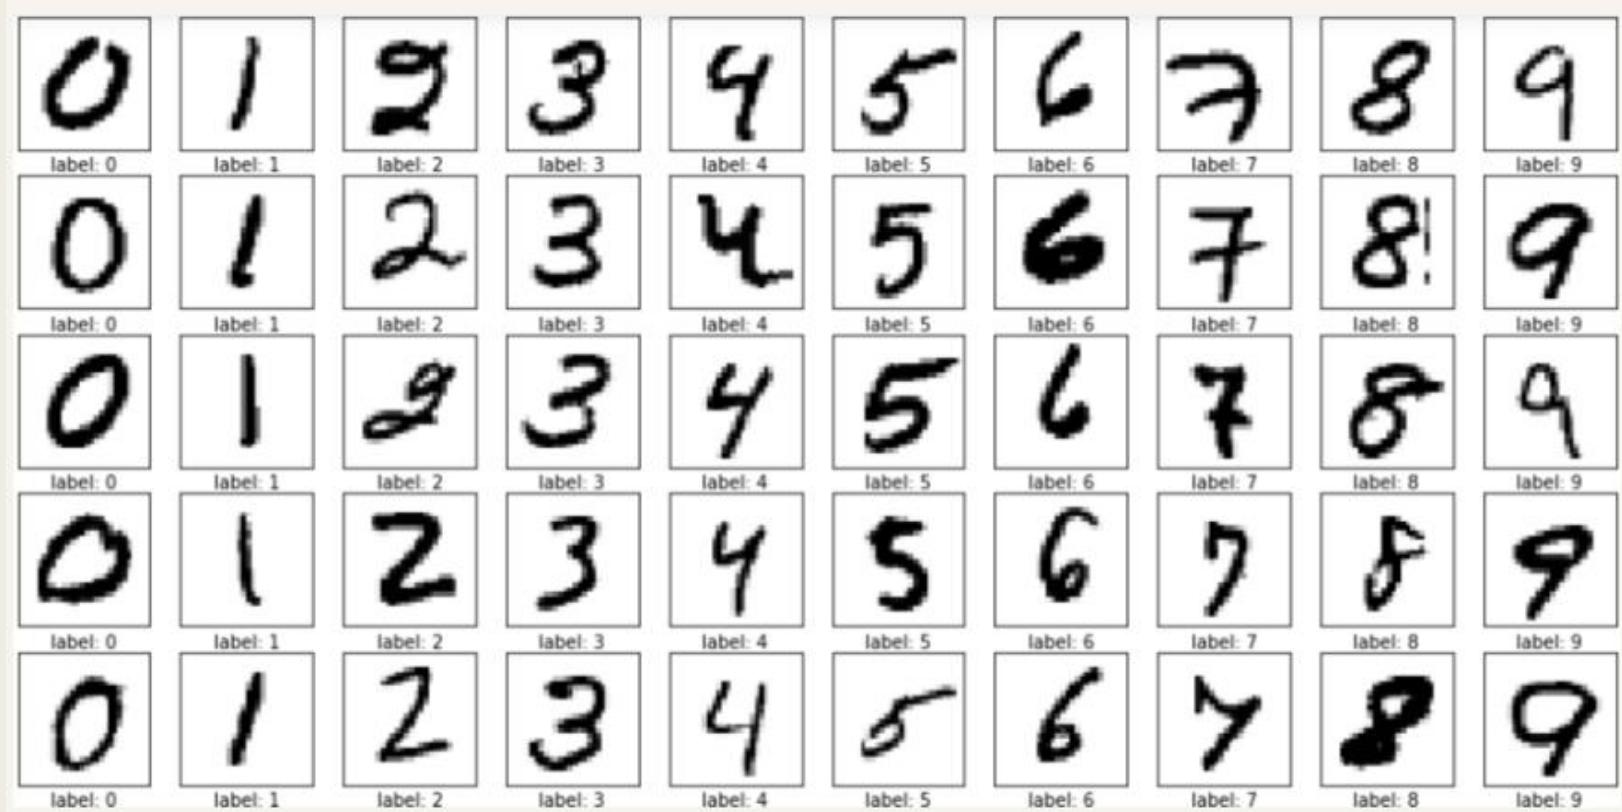
# First Steps in Deep Learning

- Choose a framework (e.g., PyTorch or TensorFlow)
- Start with a simple dataset (e.g., MNIST)
- Experiment with a small neural network
- Visualize results to understand model behavior

# Example: MNIST Dataset

- Dataset of 70,000 handwritten digits (0–9)
- Task: Classify images into correct digit
- Commonly used for benchmarking neural networks [Image: Sample MNIST digits]

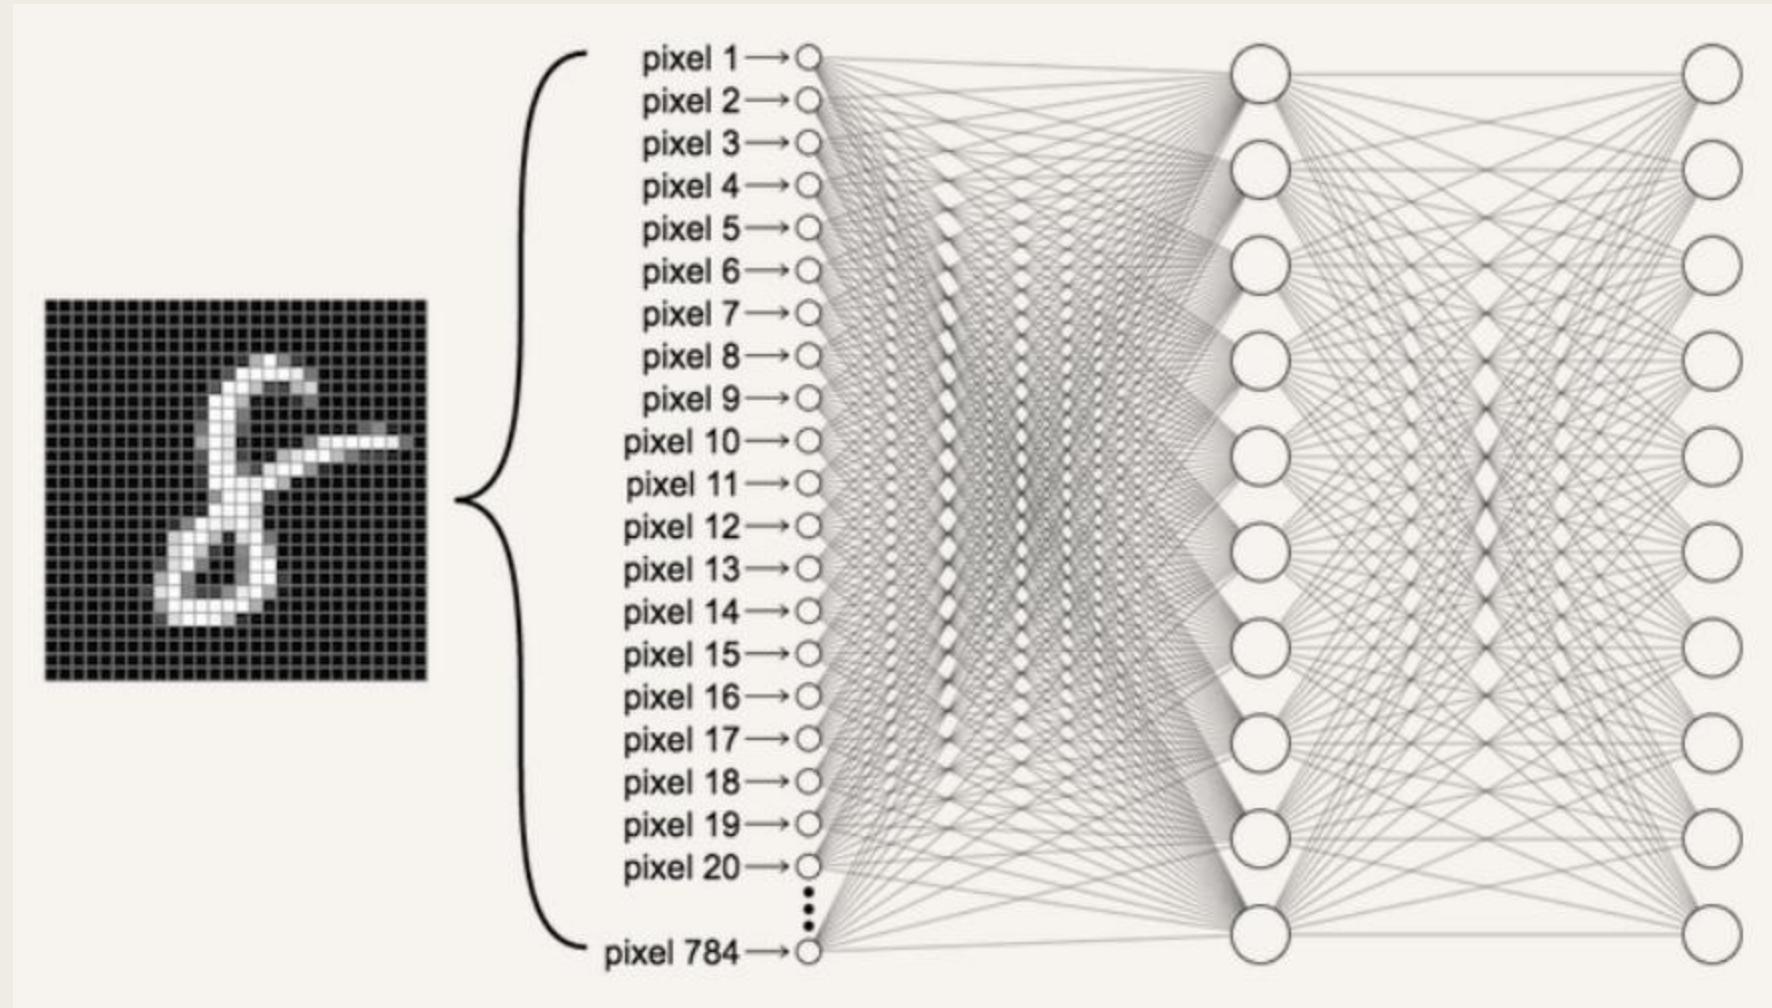
# Simple Neural Network



# Simple Neural Network

- Example: Classify handwritten digits (MNIST)
  - Layers: Dense → ReLU → Dense → Softmax
  - Output: Predicted digit (0–9)
- 
- Example coded in Keras (using tensorflow backend).
  - Input layer  $28 \times 28 = 784$
  - First layer: Dense 512, activation: relu
  - Second layer: Dense 512, activation: relu
  - Final layer: Dense 10, activation: softmax
  - Function loss: cross entropy
  - Gradient algorithm: Minibatch 128 size and RMSprop
  - Num epochs: 20

# Simple Neural Network



# Code

```
■ import tensorflow as tf
■ from tensorflow.keras import layers, models
■ # Load and preprocess MNIST dataset
■ (x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()
■ x_train, x_test = x_train / 255.0, x_test / 255.0 # Normalize pixel values to [0, 1]
■ x_train = x_train.reshape(-1, 28 * 28) # Flatten 28x28 images to 784
■ x_test = x_test.reshape(-1, 28 * 28)
■ # Build simple neural network
■ model = models.Sequential([
■     layers.Dense(128, activation='relu', input_shape=(784,)), # Dense + ReLU
■     layers.Dense(10, activation='softmax') # Dense + Softmax (10 classes)
■ ])
```

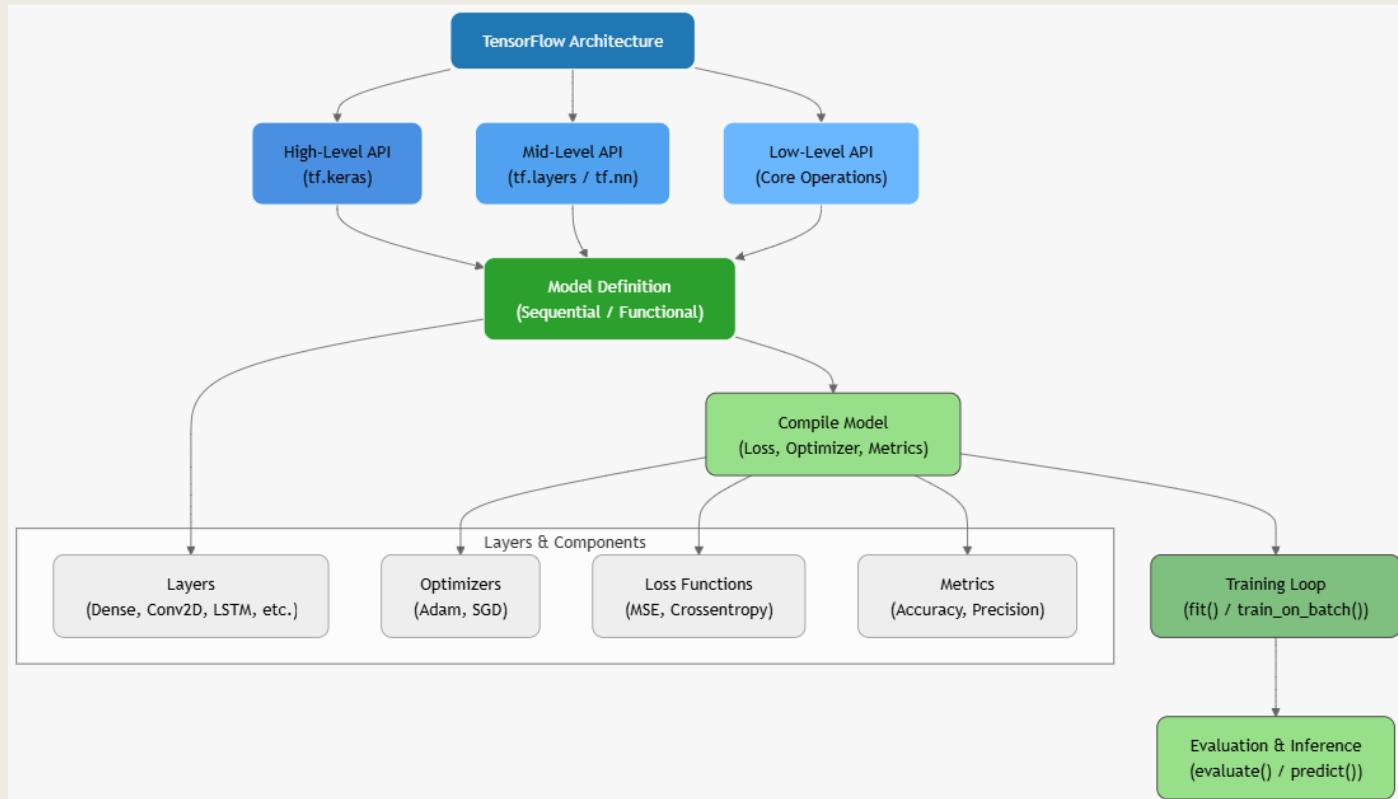
# Code

- # Compile and train
- model.compile(optimizer='adam',  
loss='sparse\_categorical\_crossentropy', metrics=['accuracy'])
- model.fit(x\_train, y\_train, epochs=5, batch\_size=32, verbose=1)
- # Evaluate on test data
- test\_loss, test\_acc = model.evaluate(x\_test, y\_test, verbose=0)
- print(f"Test accuracy: {test\_acc:.4f}")

# Importing Libraries

```
import tensorflow as tf  
from tensorflow.keras import layers, models
```

- **tensorflow** – main deep learning framework.
- **layers** – defines neural network layers (Dense, Conv, etc.).
- **models** – builds and manages neural network architectures.



# Loading the MNIST Dataset

```
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()
```

- MNIST = handwritten digits (0–9) dataset.
- x\_train → images for training.
- y\_train → labels for training.
- x\_test, y\_test → images and labels for testing.
- Dataset size: 60,000 training + 10,000 testing samples.

# Normalizing Pixel Values

```
x_train, x_test = x_train / 255.0, x_test / 255.0
```

- Converts pixel range from 0-255 → 0-1.
- Helps model train faster and perform better.
- Normalization ensures **numerical stability** and smoother gradients.

# Flattening Images

```
x_train = x_train.reshape(-1, 28 * 28)  
x_test = x_test.reshape(-1, 28 * 28)
```

- Each image =  $28 \times 28$  pixels → 784 features.
- Converts 2D image into a 1D vector for input to the dense layers.
- -1 lets NumPy automatically calculate batch size.

# Building the Neural Network

```
model = models.Sequential([
    layers.Dense(128, activation='relu', input_shape=(784,)),
    layers.Dense(10, activation='softmax')
])
```

- **Sequential model:** stack of layers in order.
- **Layer 1:** 128 neurons, ReLU activation → learns features.
- **Layer 2:** 10 neurons, Softmax → outputs class probabilities (digits 0–9).

# Compiling the Model

```
model.compile(  
    optimizer='adam',  
    loss='sparse_categorical_crossentropy',  
    metrics=['accuracy'])
```

- **Optimizer:** Adam (adaptive learning rate).
- **Loss:** Sparse categorical cross-entropy (for integer labels).
- **Metric:** Accuracy (percentage of correct predictions).

# Training the Model

```
model.fit(x_train, y_train, epochs=5, batch_size=32, verbose=1)
```

- **epochs=5**: train over full dataset 5 times.
- **batch\_size=32**: updates weights every 32 samples.
- **verbose=1**: shows progress (e.g., 1500/1500 steps).

# Evaluating the Model

```
test_loss, test_acc = model.evaluate(x_test, y_test, verbose=0)
print(f"Test accuracy: {test_acc:.4f}")
```

- Uses unseen test data.
- Returns **test loss** and **test accuracy**.
- Accuracy typically ~0.97–0.98 (97–98%).

# Summary

- MNIST is a great dataset to start with.
- Neural networks learn from features → outputs probabilities.
- Proper normalization, model structure, and tuning are essential.
- Achieved ~98% accuracy with simple architecture.